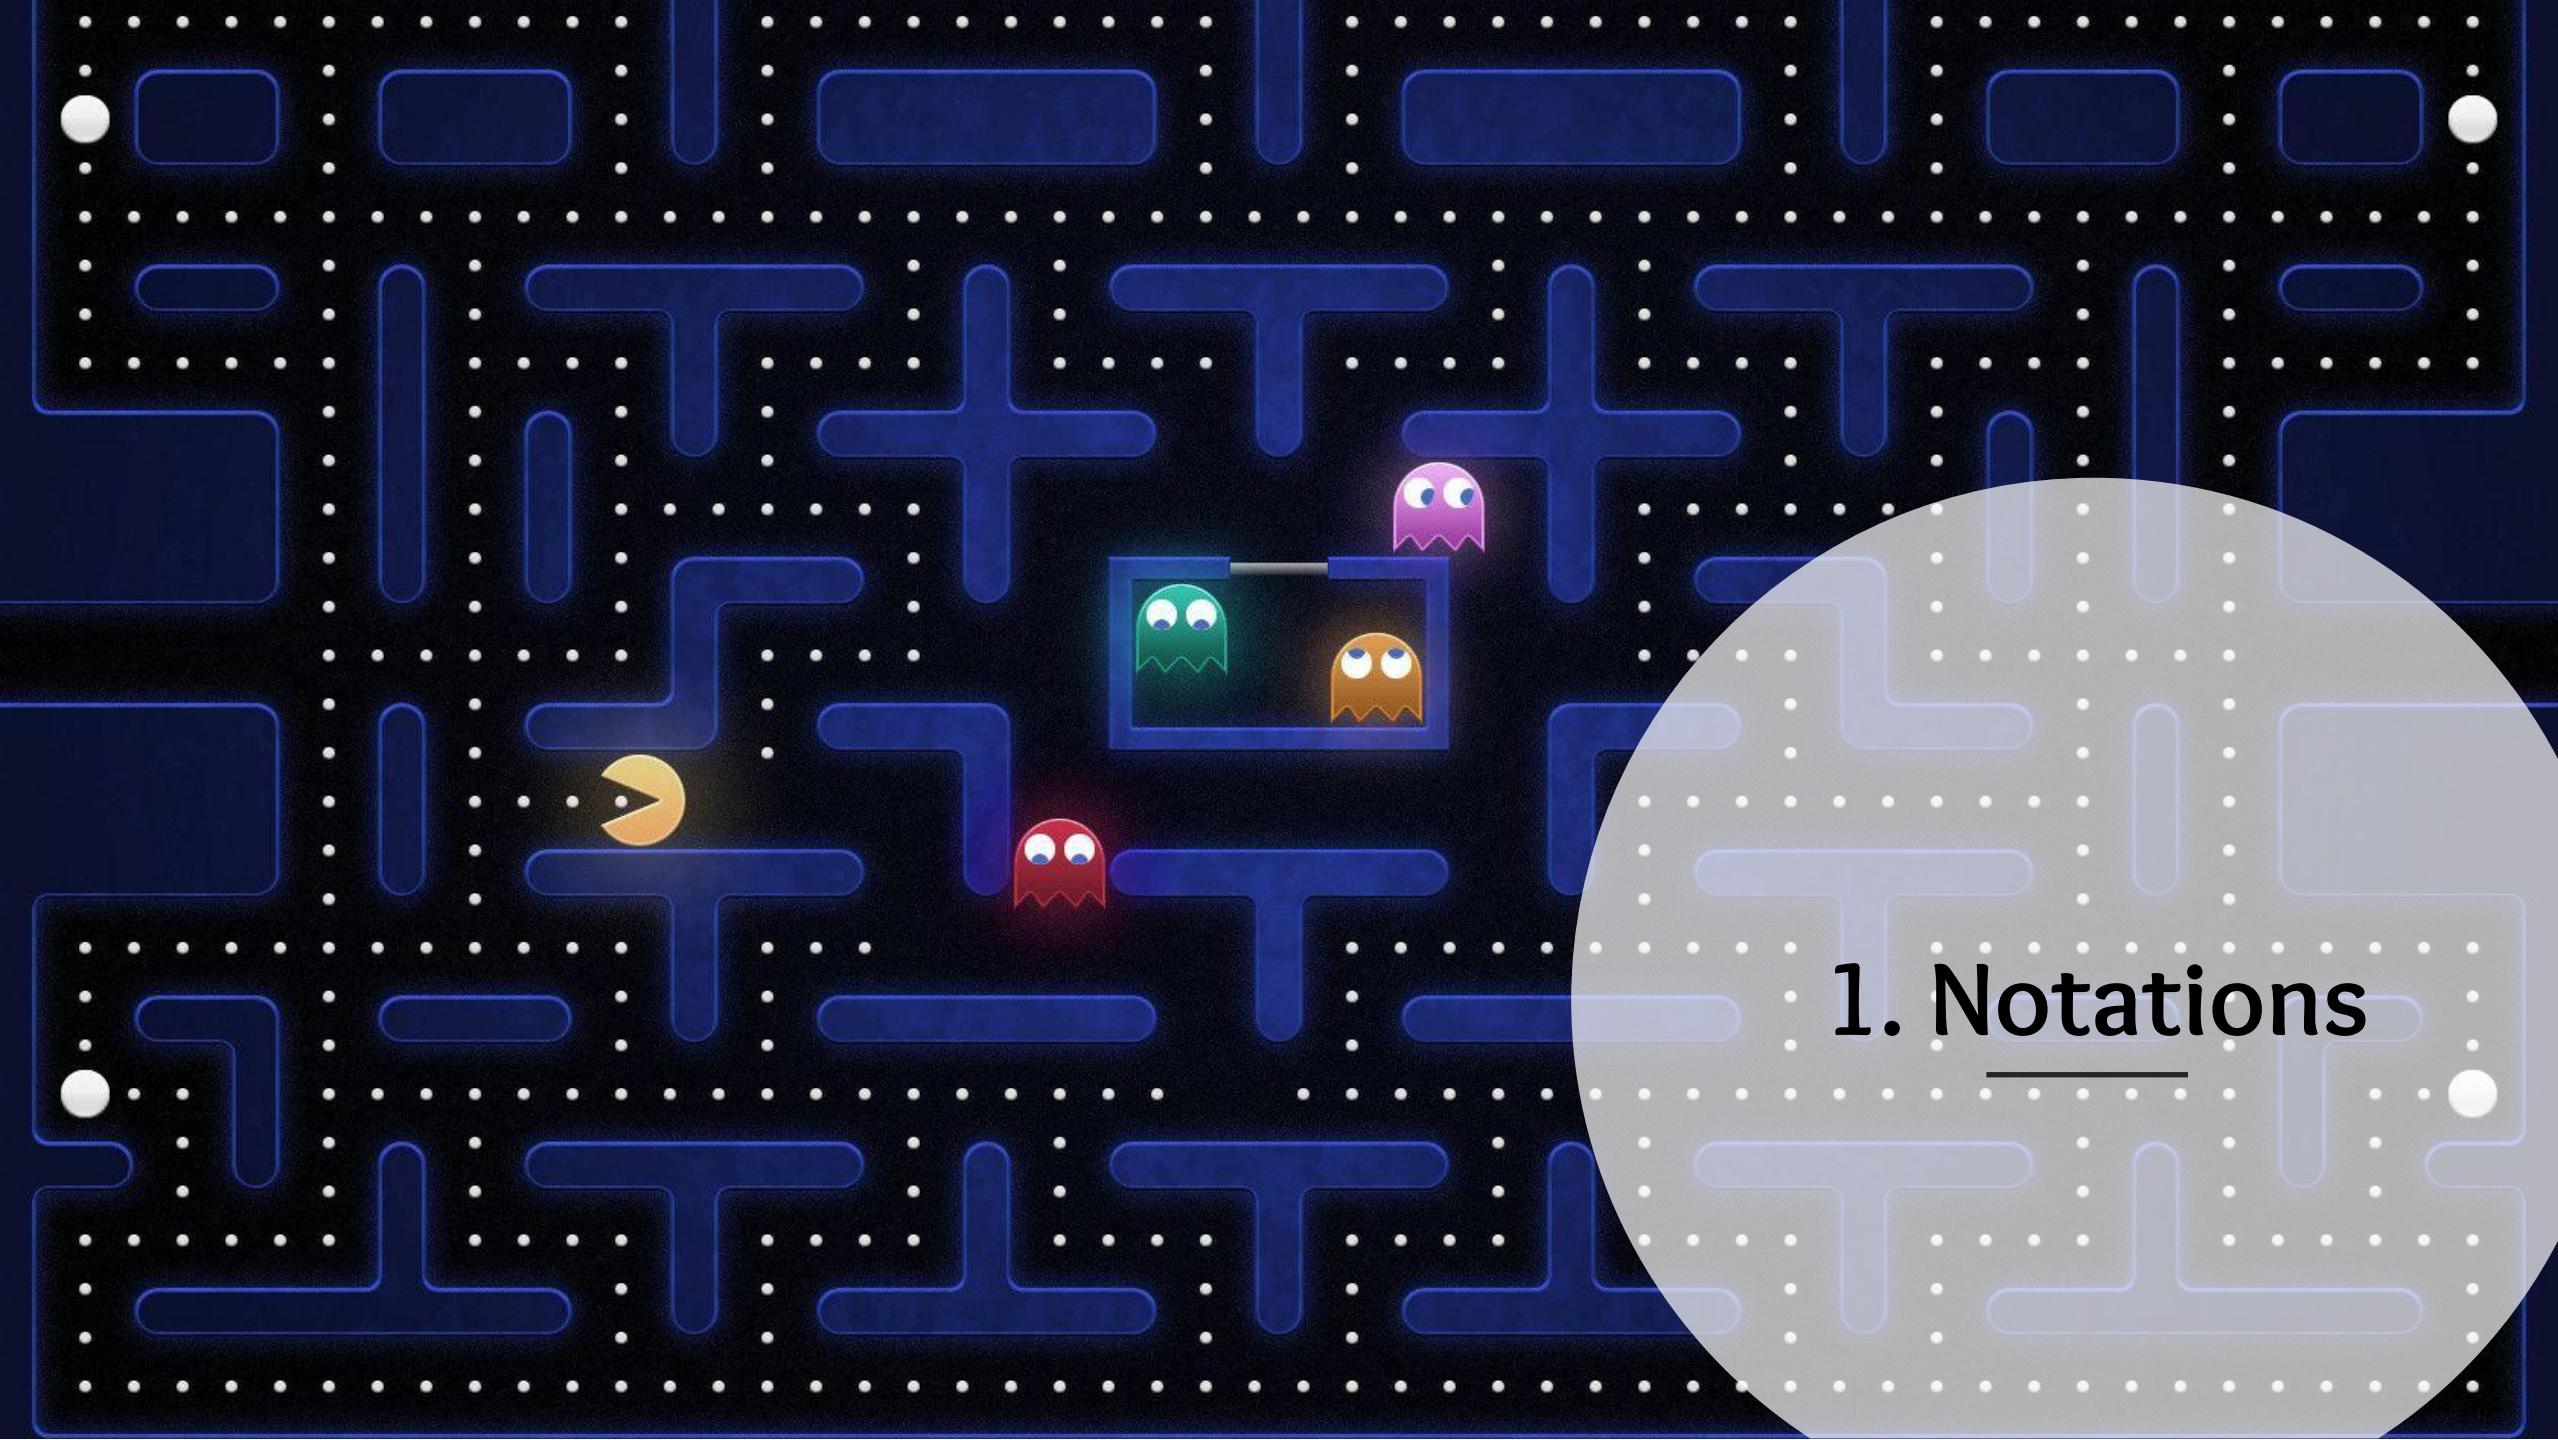
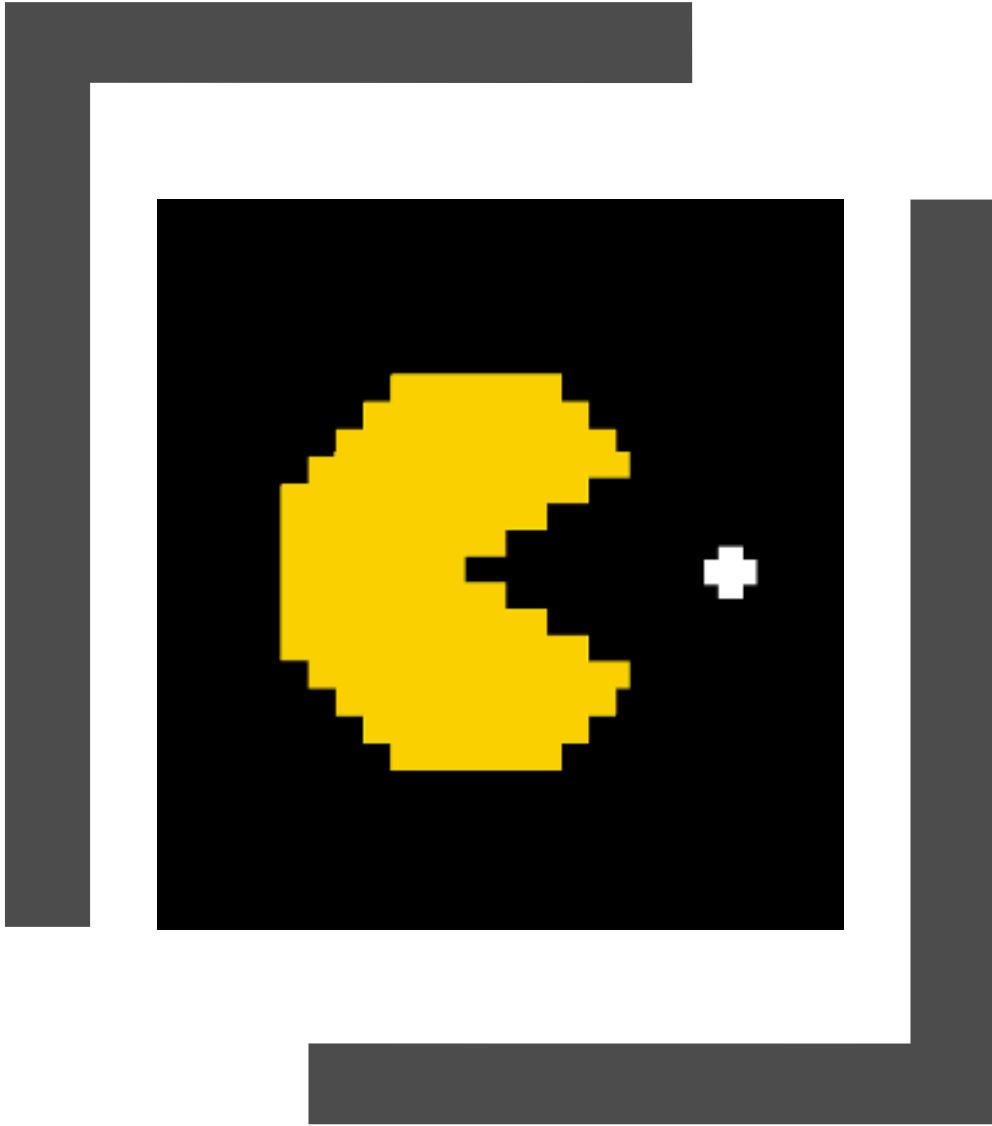


CS 3600 Project 1 Review

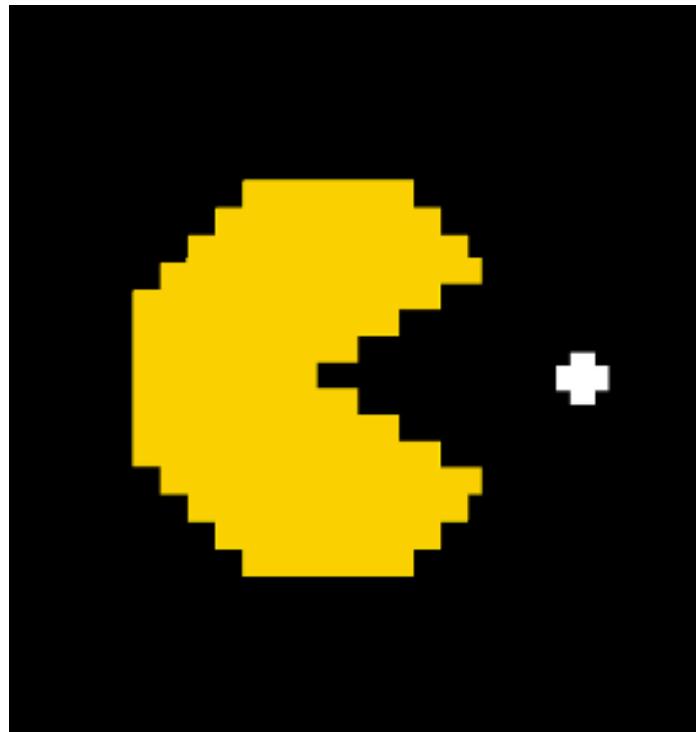


1. Notations



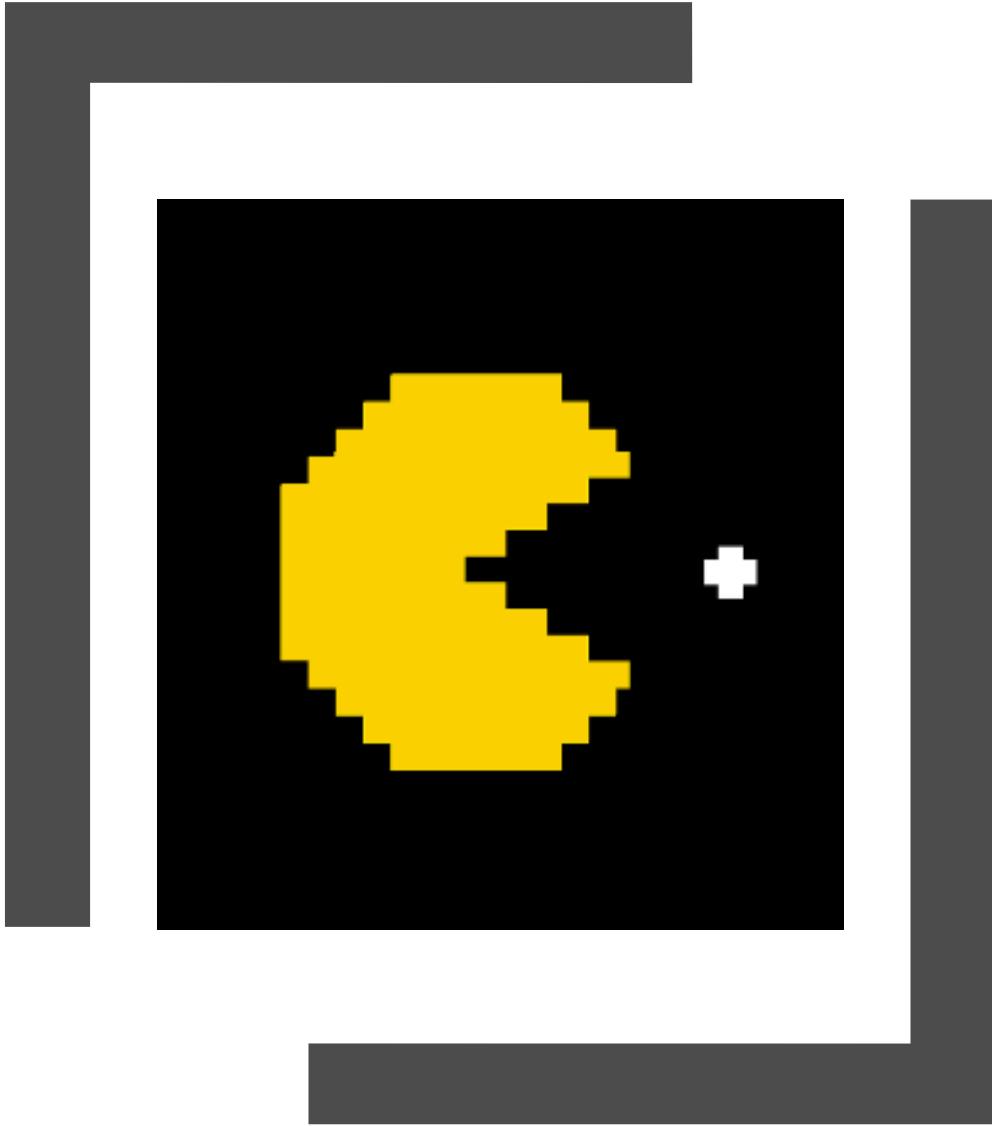
Notations: State

- Representation of where Pacman is on the map!
- Represented as *Euclidean* coordinates, or just a label (e.g. "A", "B", ⋯)
- Will refine notion of "state" in later question of HW



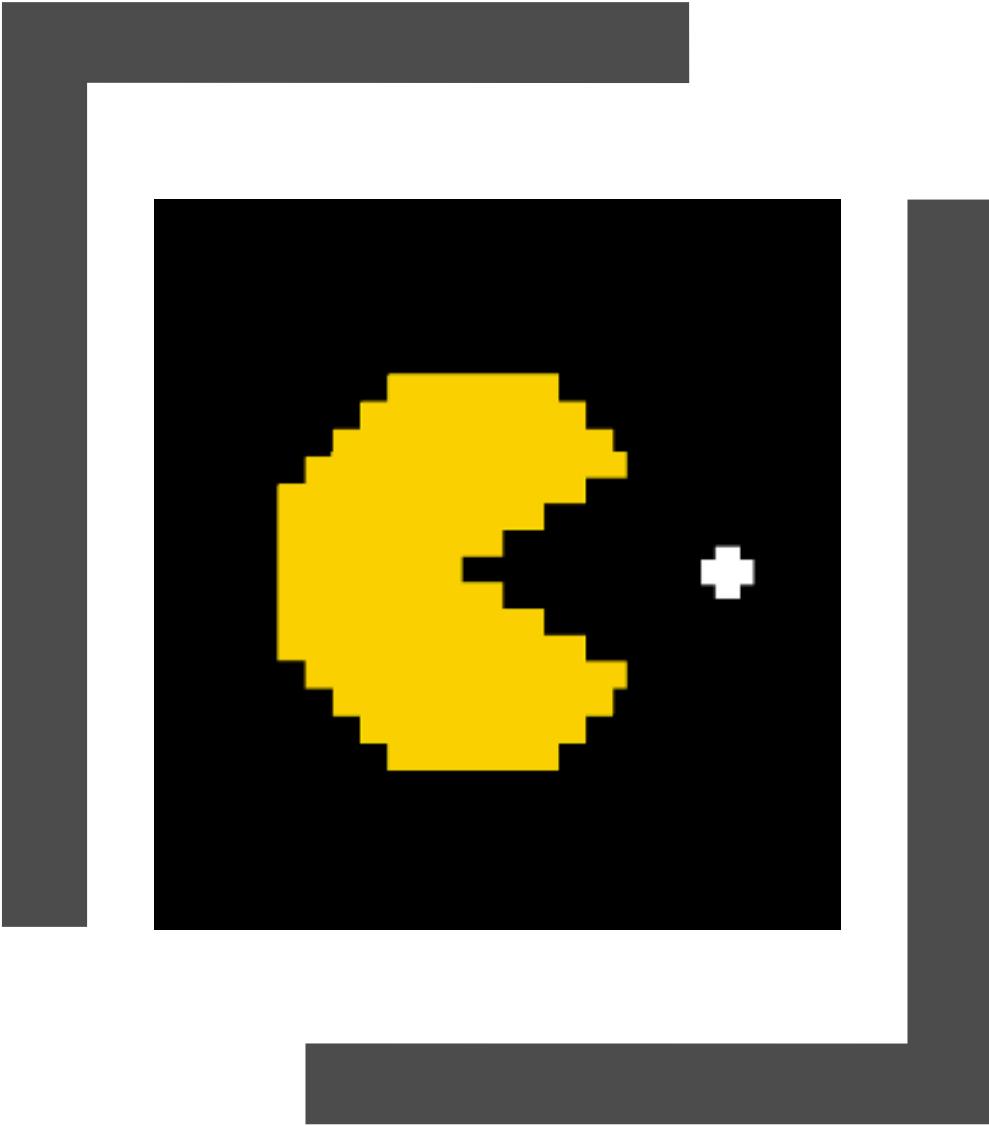
Notations: Successors

- The next possible states
- Tuple containing *three* items:
 - Current state
 - Action taken to reach it
 - Cost of action
- Call `getSuccessors()`



Notations: Path

- Sequences of actions returned from your search functions
- How to get from *start state S* to *goal state G*



Notations: Goal

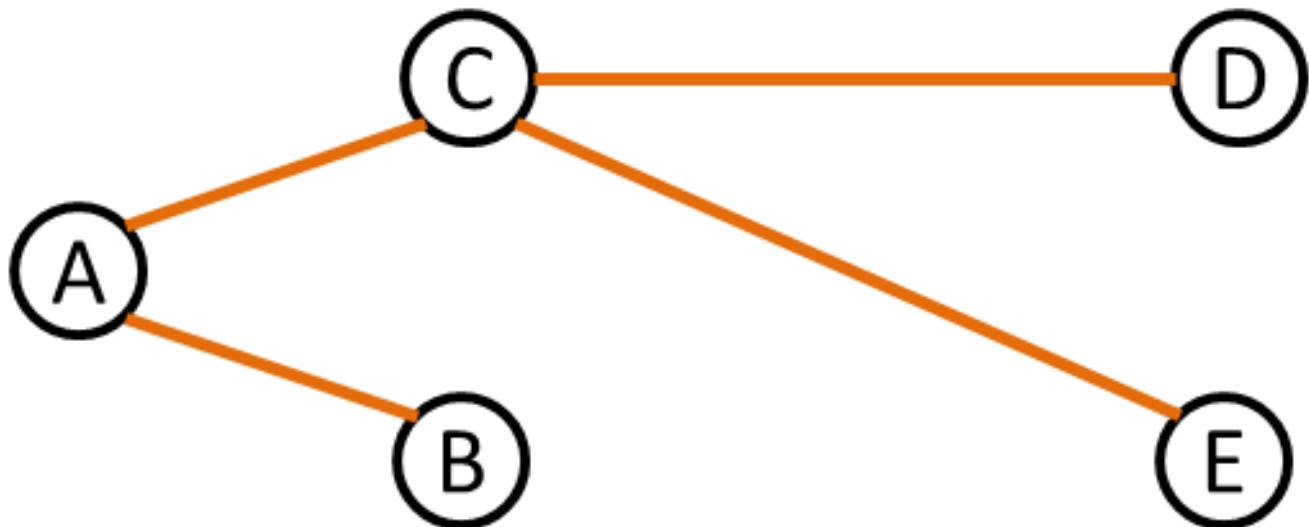
- State that you want to reach
- Final destination
- Check using `isGoalState()`



2. Graph Search Algorithms

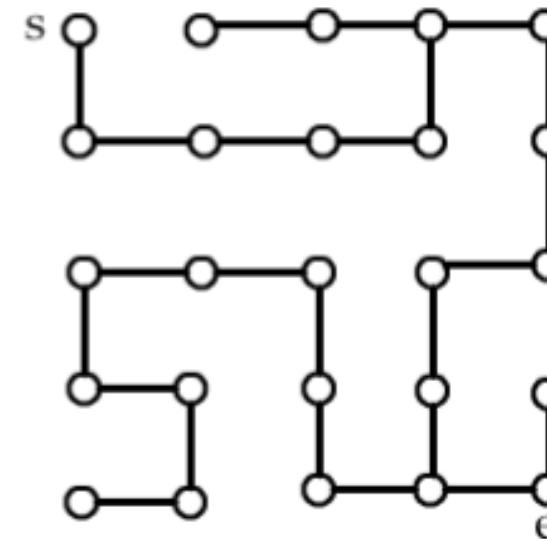
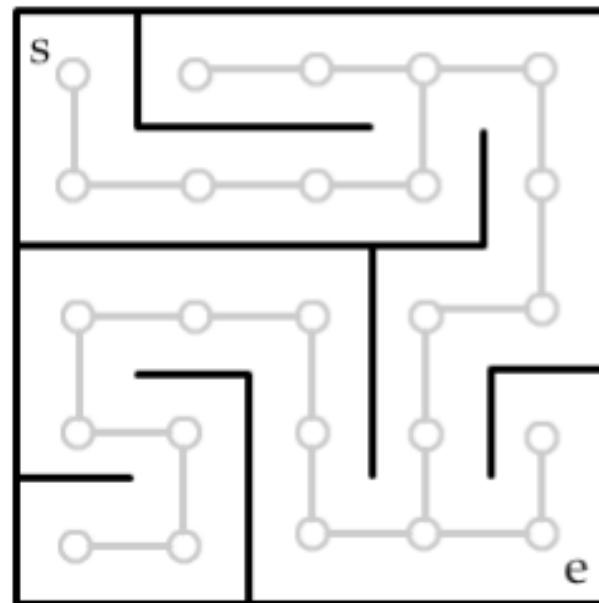
Graphs

- Graphs have **nodes** and **edges**



Representing grids as a graph

- Each cell becomes a node. Edges connect to adjacent cells.



Graph Searching Algorithms

- Begin at the start node and keep searching till we find the goal node
- Different algorithms → Different orders in which to search the nodes of the graph
- Open-list contains nodes that we've seen but haven't explored yet. Each iteration, we take a node off the open-list, and add its neighbors to the open-list.

Generic Search Algorithm

Algorithm 1 Generic Graph Search Algorithm

Input: *startState*, the state at which we begin our search.

Output: *path*, the sequence of actions to get from *startState* to a goal state

```
1: // Initialize variables
2: openList  $\leftarrow$  dataStructure()       $\triangleright$  dataStructure  $\in \{\text{Stack, Queue, PriorityQueue}\}
3: closedList  $\leftarrow$  set()                   $\triangleright$  Only explored nodes will go in here
4: metadata  $\leftarrow$  createMetadata(startState, nil)
5: current  $\leftarrow$  (startState, metadata)       $\triangleright$  Define current as the state plus metadata
6:
7: // Search until goal state found.
8: while current.state is not a goal state do           $\triangleright$  Continue until goal reached
9:   p  $\leftarrow$  current.state                 $\triangleright$  Define p as the state we're interested in, the parent
10:  if p is not in closedList then            $\triangleright$  If p is unexplored, explore it
11:    closedList.add(p)                       $\triangleright$  Mark p as explored.
12:    for s  $\in$  successors(p) do            $\triangleright$  Retrieve all successors for the parent
13:      s.metadata  $\leftarrow$  createMetadata(s, p)   $\triangleright$  Create metadata using both p and s
14:      openList.add((s.state, s.metadata))   $\triangleright$  Add the successor onto the open list
15:    current  $\leftarrow$  openList.next()           $\triangleright$  Fetch a new current from the openList
16:
17: // Reconstruct path from startState to goal state
18: path  $\leftarrow$  reconstructPath()
19: return path$ 
```

Data Structures

Same code for all! Just need to change the open-list:

- **Queue** → Breadth-First Search (BFS)
- **Stack** → Depth-First Search (DFS)
- **Priority Queue** → Uniform Cost Search (UCS) and A*
(Note: The priority functions are different for these two algorithms)

Metadata

We might need some metadata (~additional information) that we need to keep track of.

Wish to accomplish 2 things mainly:

1. Help keep track of the path taken to get to the goal
2. Allow priority of state to factor in for UCS and A*

Note:

- The metadata is problem-dependent
- Passing metadata in python can be easy – put everything in a tuple (including state), and store that tuple in the open-list.

1. Keeping track of path to goal

- **Full-Length Sequence of Actions:**
 - Simply keep a full-length list of actions taken to reach current state
 - So for each successor, the metadata will contain an entire sequence of actions to reach the parent, and a new action will be added to the end of that sequence to signify the action that takes you from s to that successor
 - *Once you complete the graph search, you'll have the full sequence of actions to reach the goal in current.metadata*
- **Parent Hash-map**
 - Note the action was taken to reach the successor, and we add this into the metadata.
 - When we begin exploring a state , we need to mark in a hashmap (*dict* in Python) the action taken to reach this state.
 - *Once we reach the goal state, we use the hashmap to back-track our way through the actions taken from each parent, until we reach the startState*
- **Linked List**

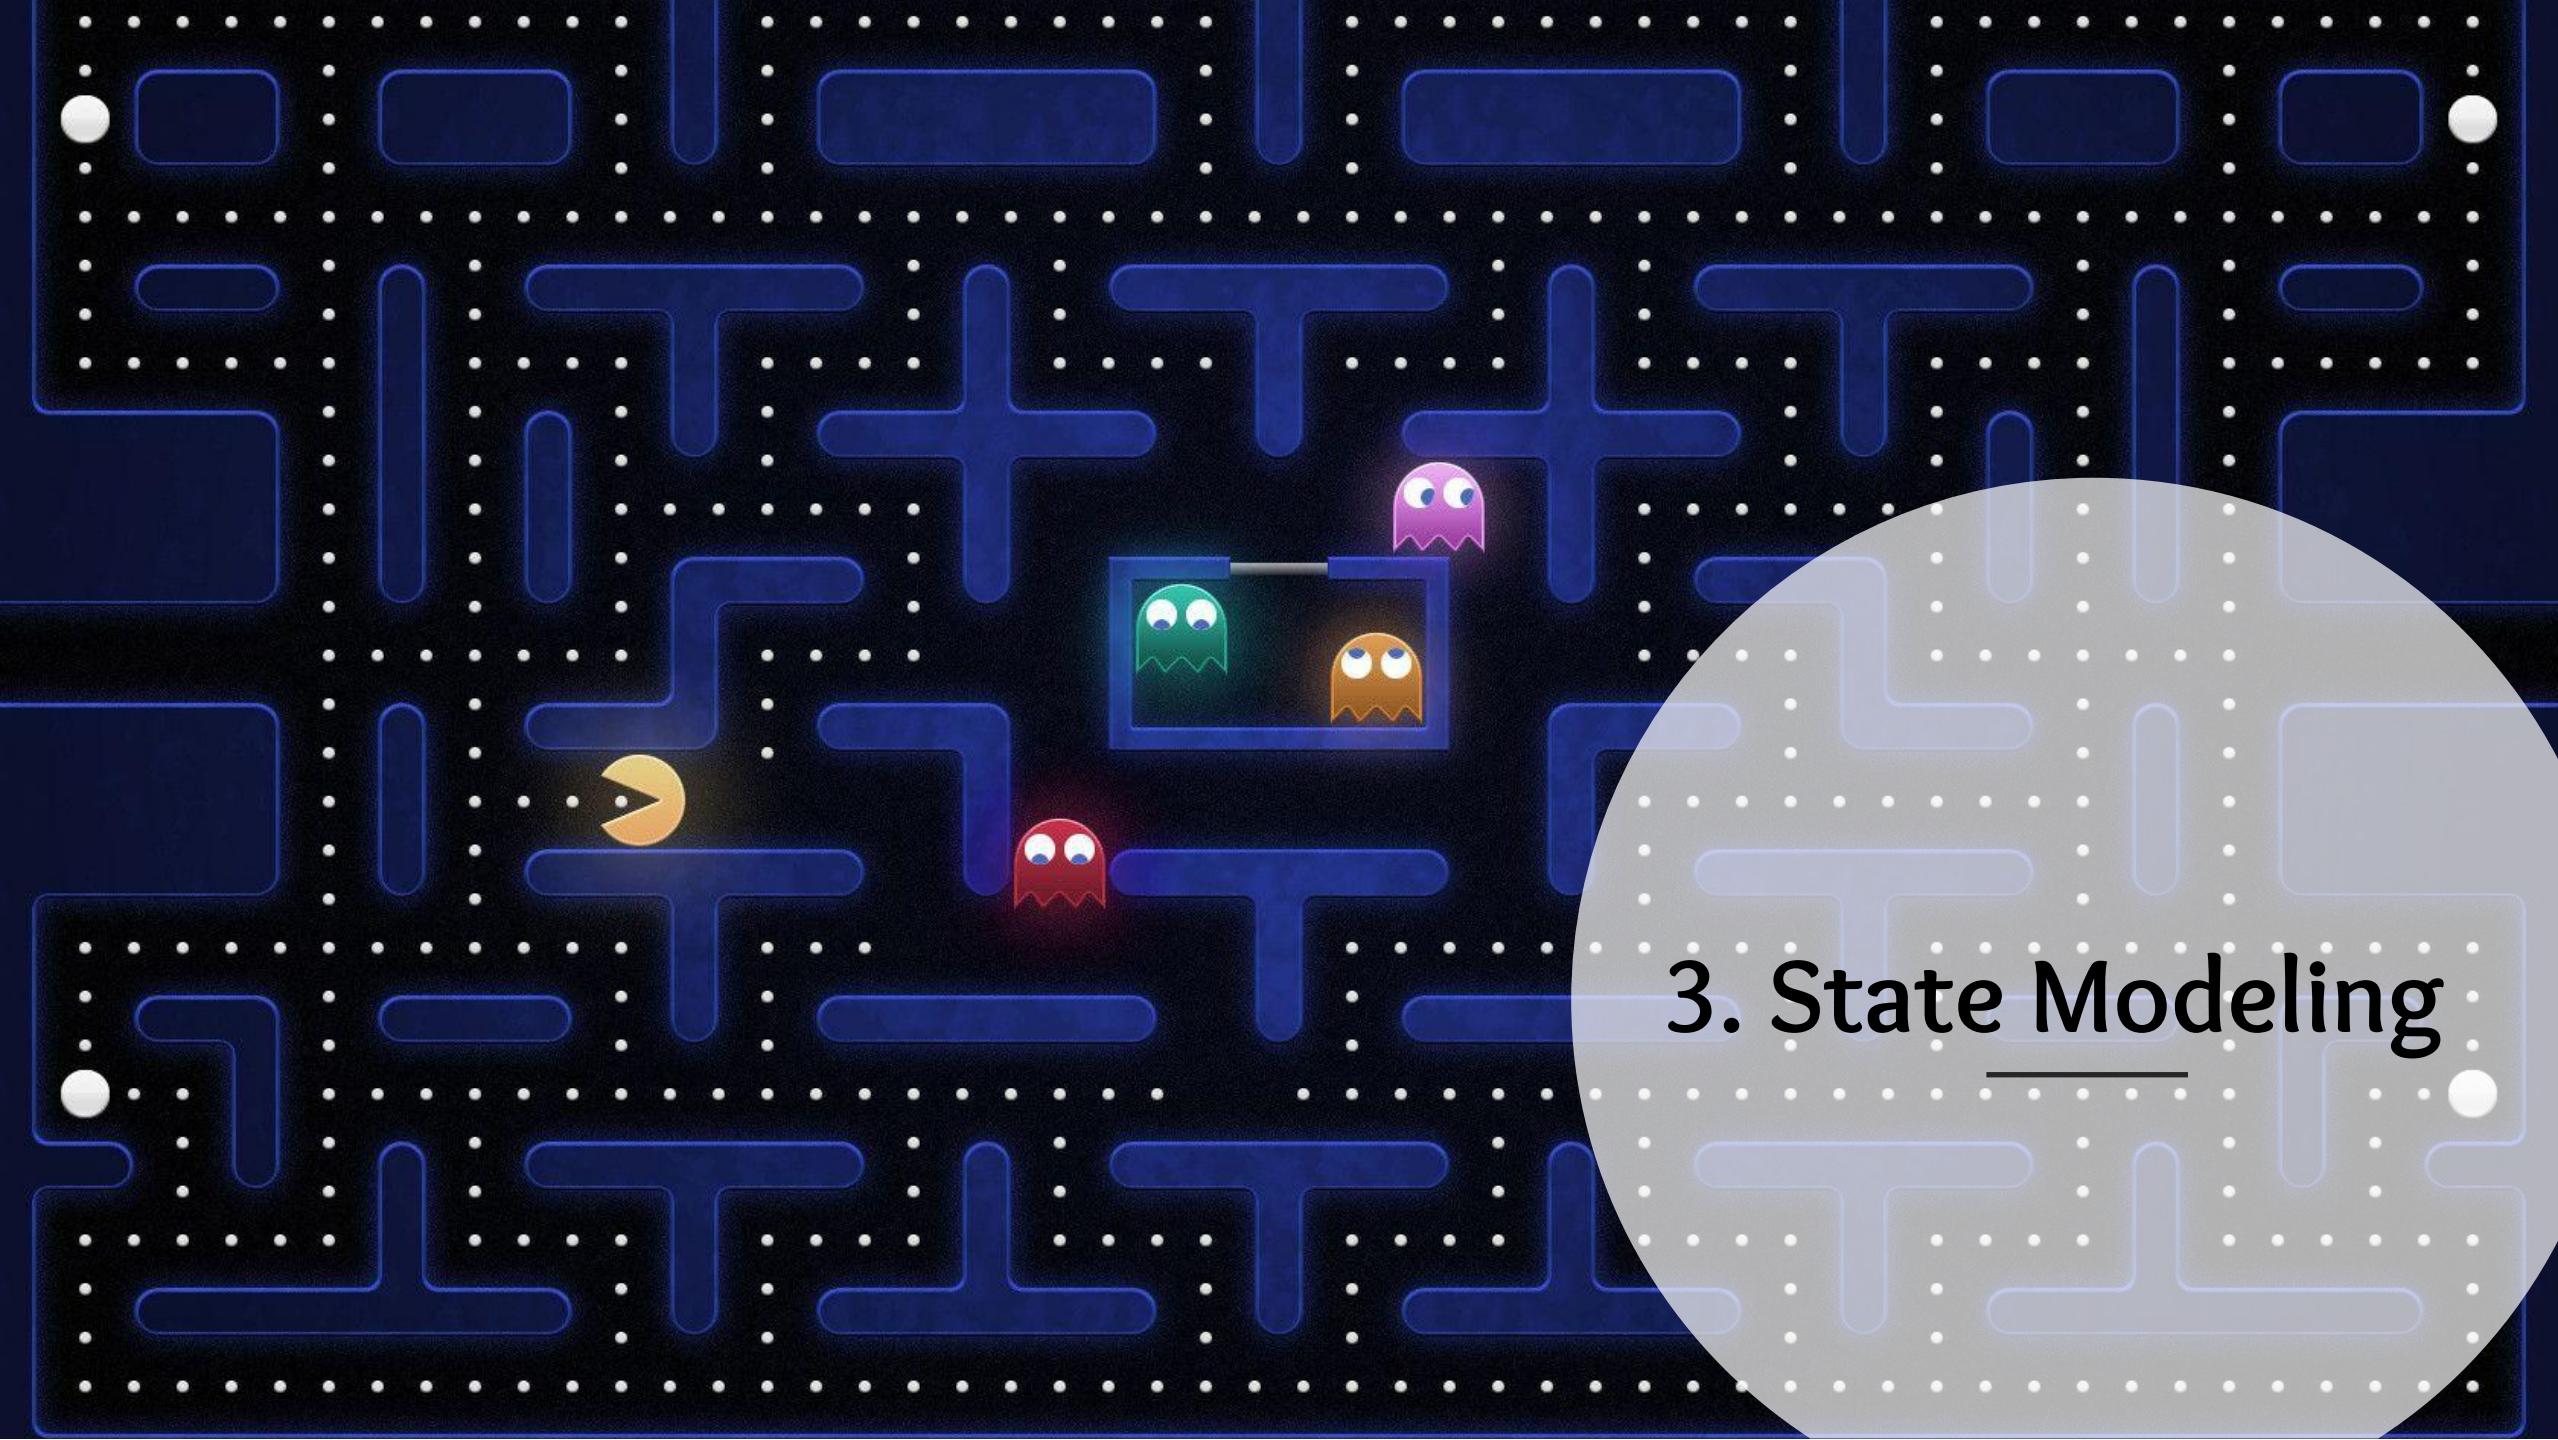
2. Determining Priority for A* and UCS

Priority needs to be passed around in the *metadata*.

- **UCS:** priority(s) = $g(s)$
- **A***: priority(s) = $g(s) + h(s)$

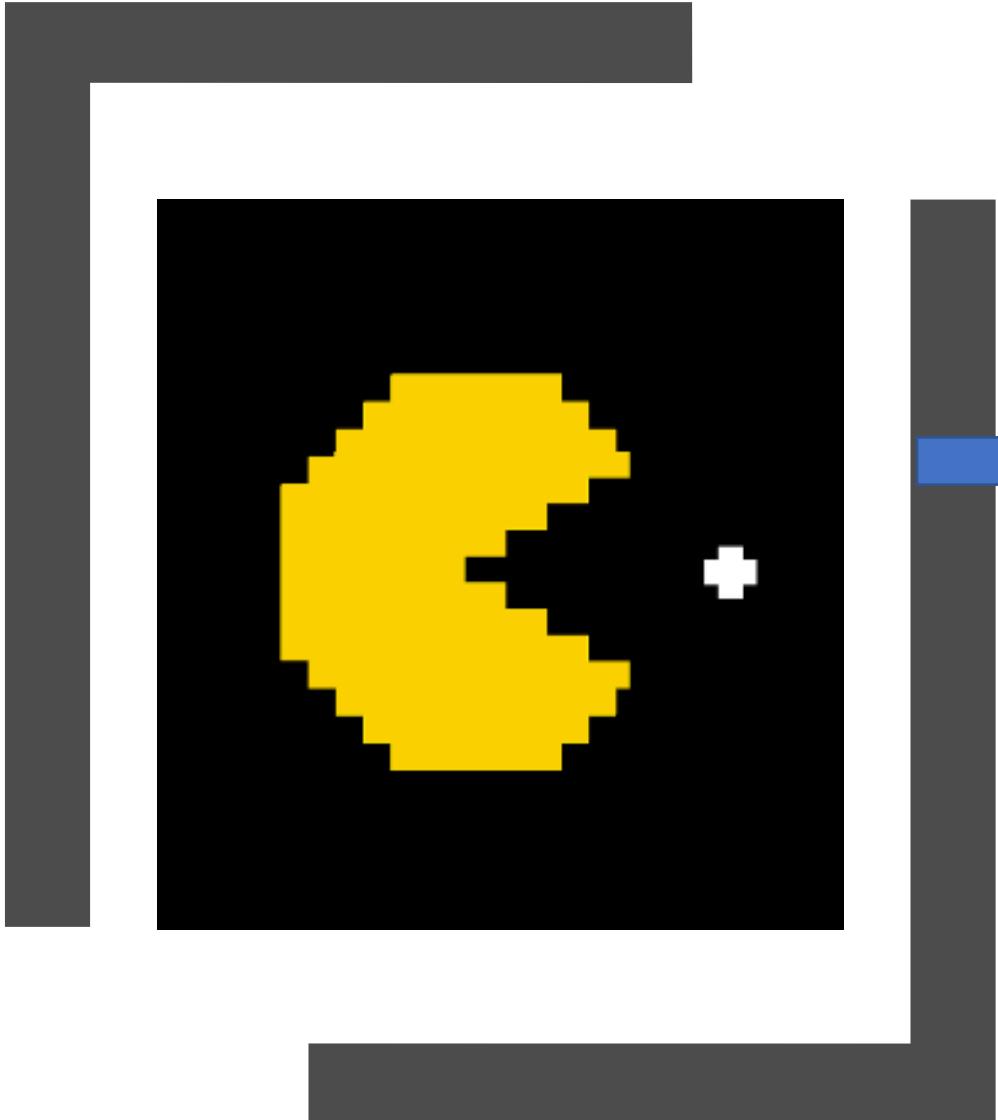
$g(s)$ tells us the cost of reaching state s

$h(s)$ approximates the remaining cost of reaching a goal state, starting from the current state s

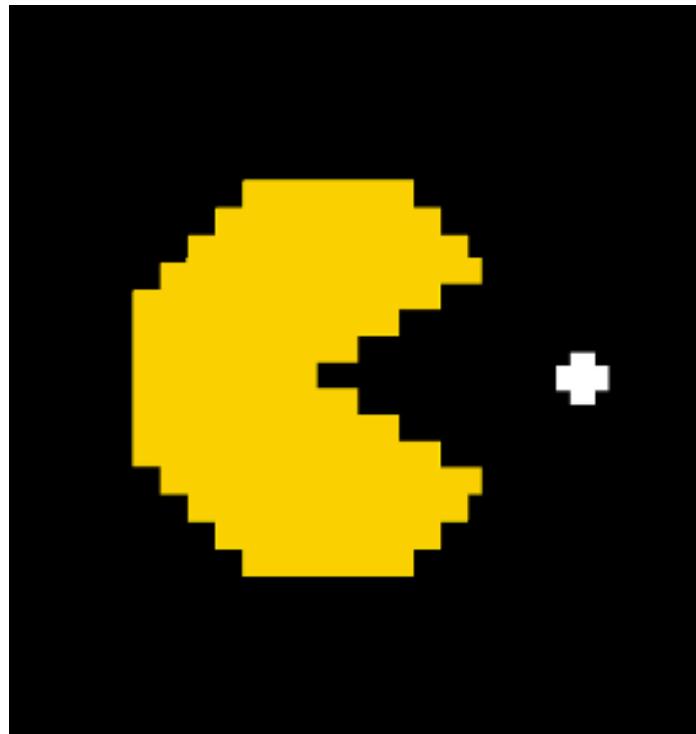
A black and white illustration of a Pac-Man maze. The maze is composed of blue T-shaped and U-shaped paths on a dark background. Four small, semi-transparent ghost characters are scattered throughout the maze: a pink ghost at the top right, a green ghost in the center, an orange ghost below it, and a red ghost at the bottom left. A single glowing yellow circle, representing a power pellet, is located near the bottom left ghost.

3. State Modeling

How do we model a search problem?

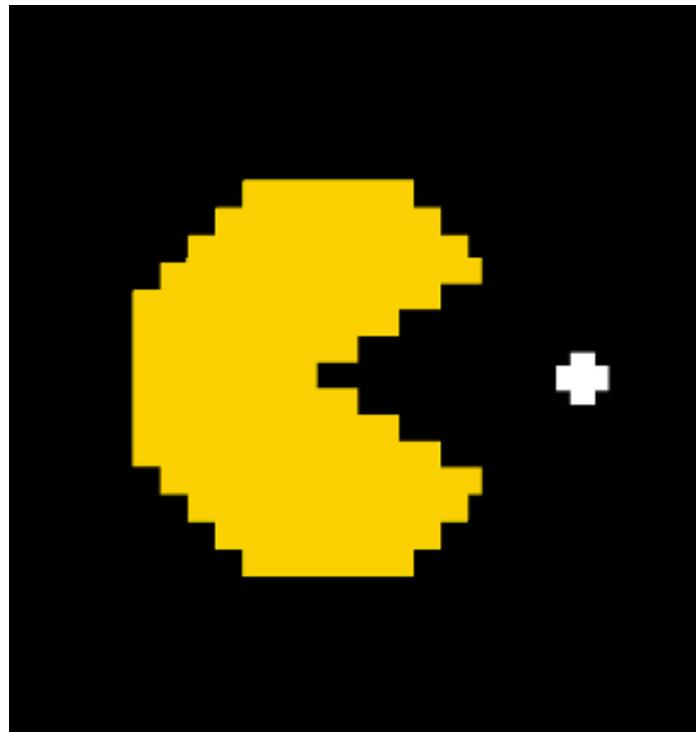


- States
- Successor Function
- Goal Function



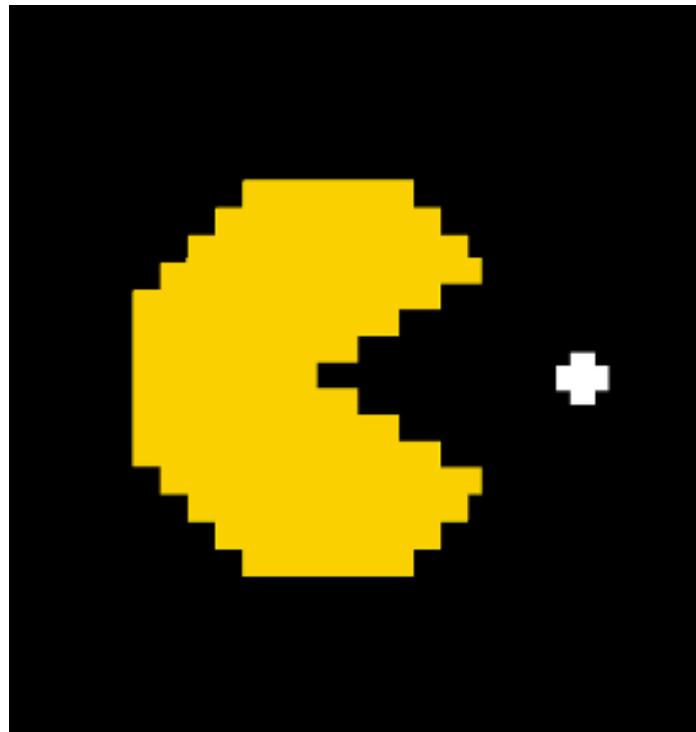
States: Encoding Matters

- Only model what you need to solve the problem
- What do we need to know about each state?
 - Its successors
 - If it is a goal state
- Don't try to model everything



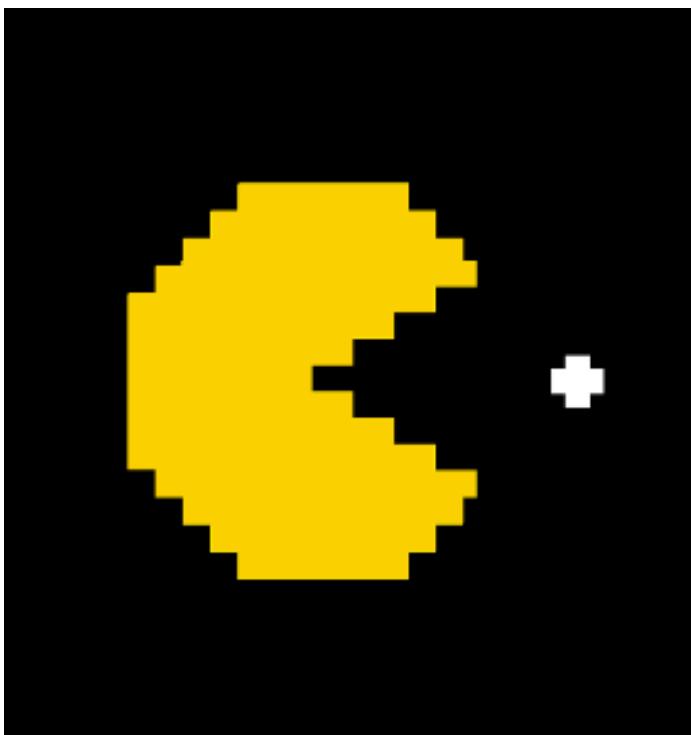
How do we model a search problem?

- States
- Successor Function
- Goal Function



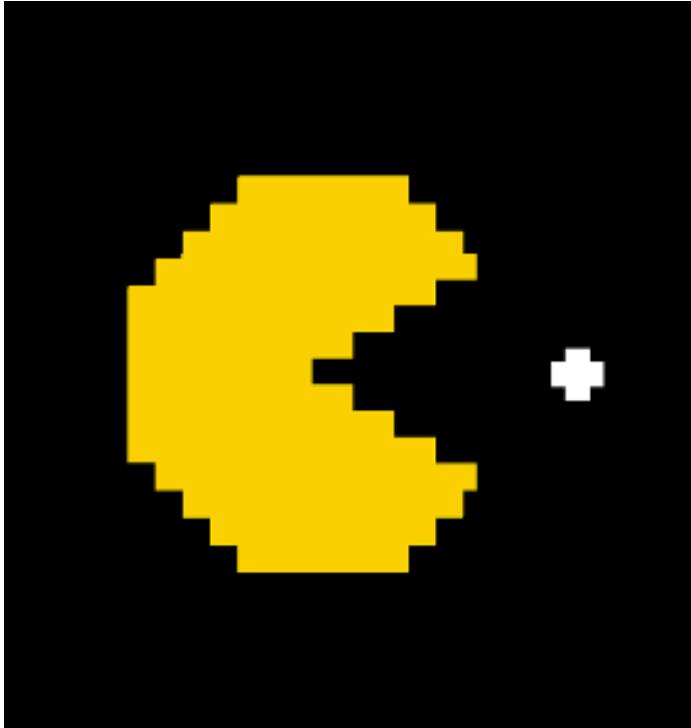
Successor Function: Efficiency Matters

- How often is our successor function called?
 - Often, every time we add nodes to our open list
- Don't make this function slow to compute



How do we model a search problem?

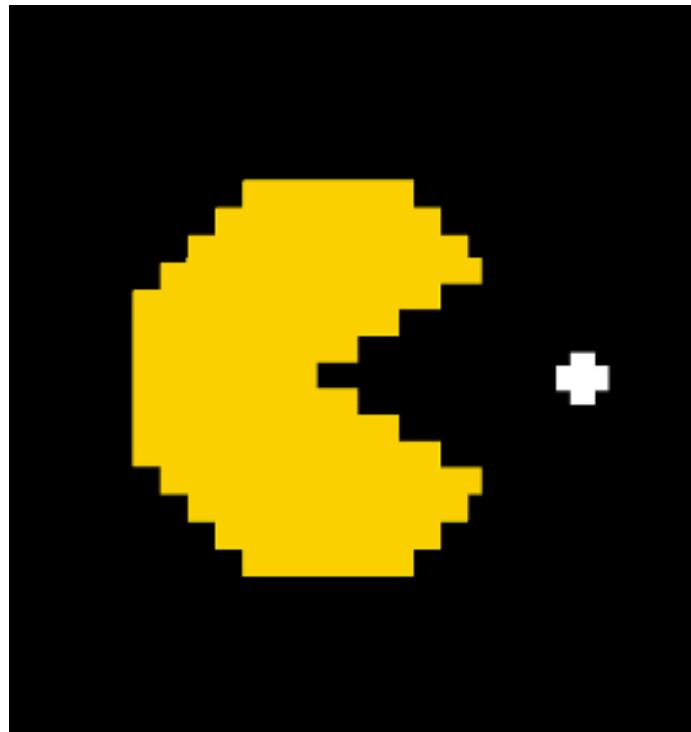
- States
- Successor Function
- **Goal Function**



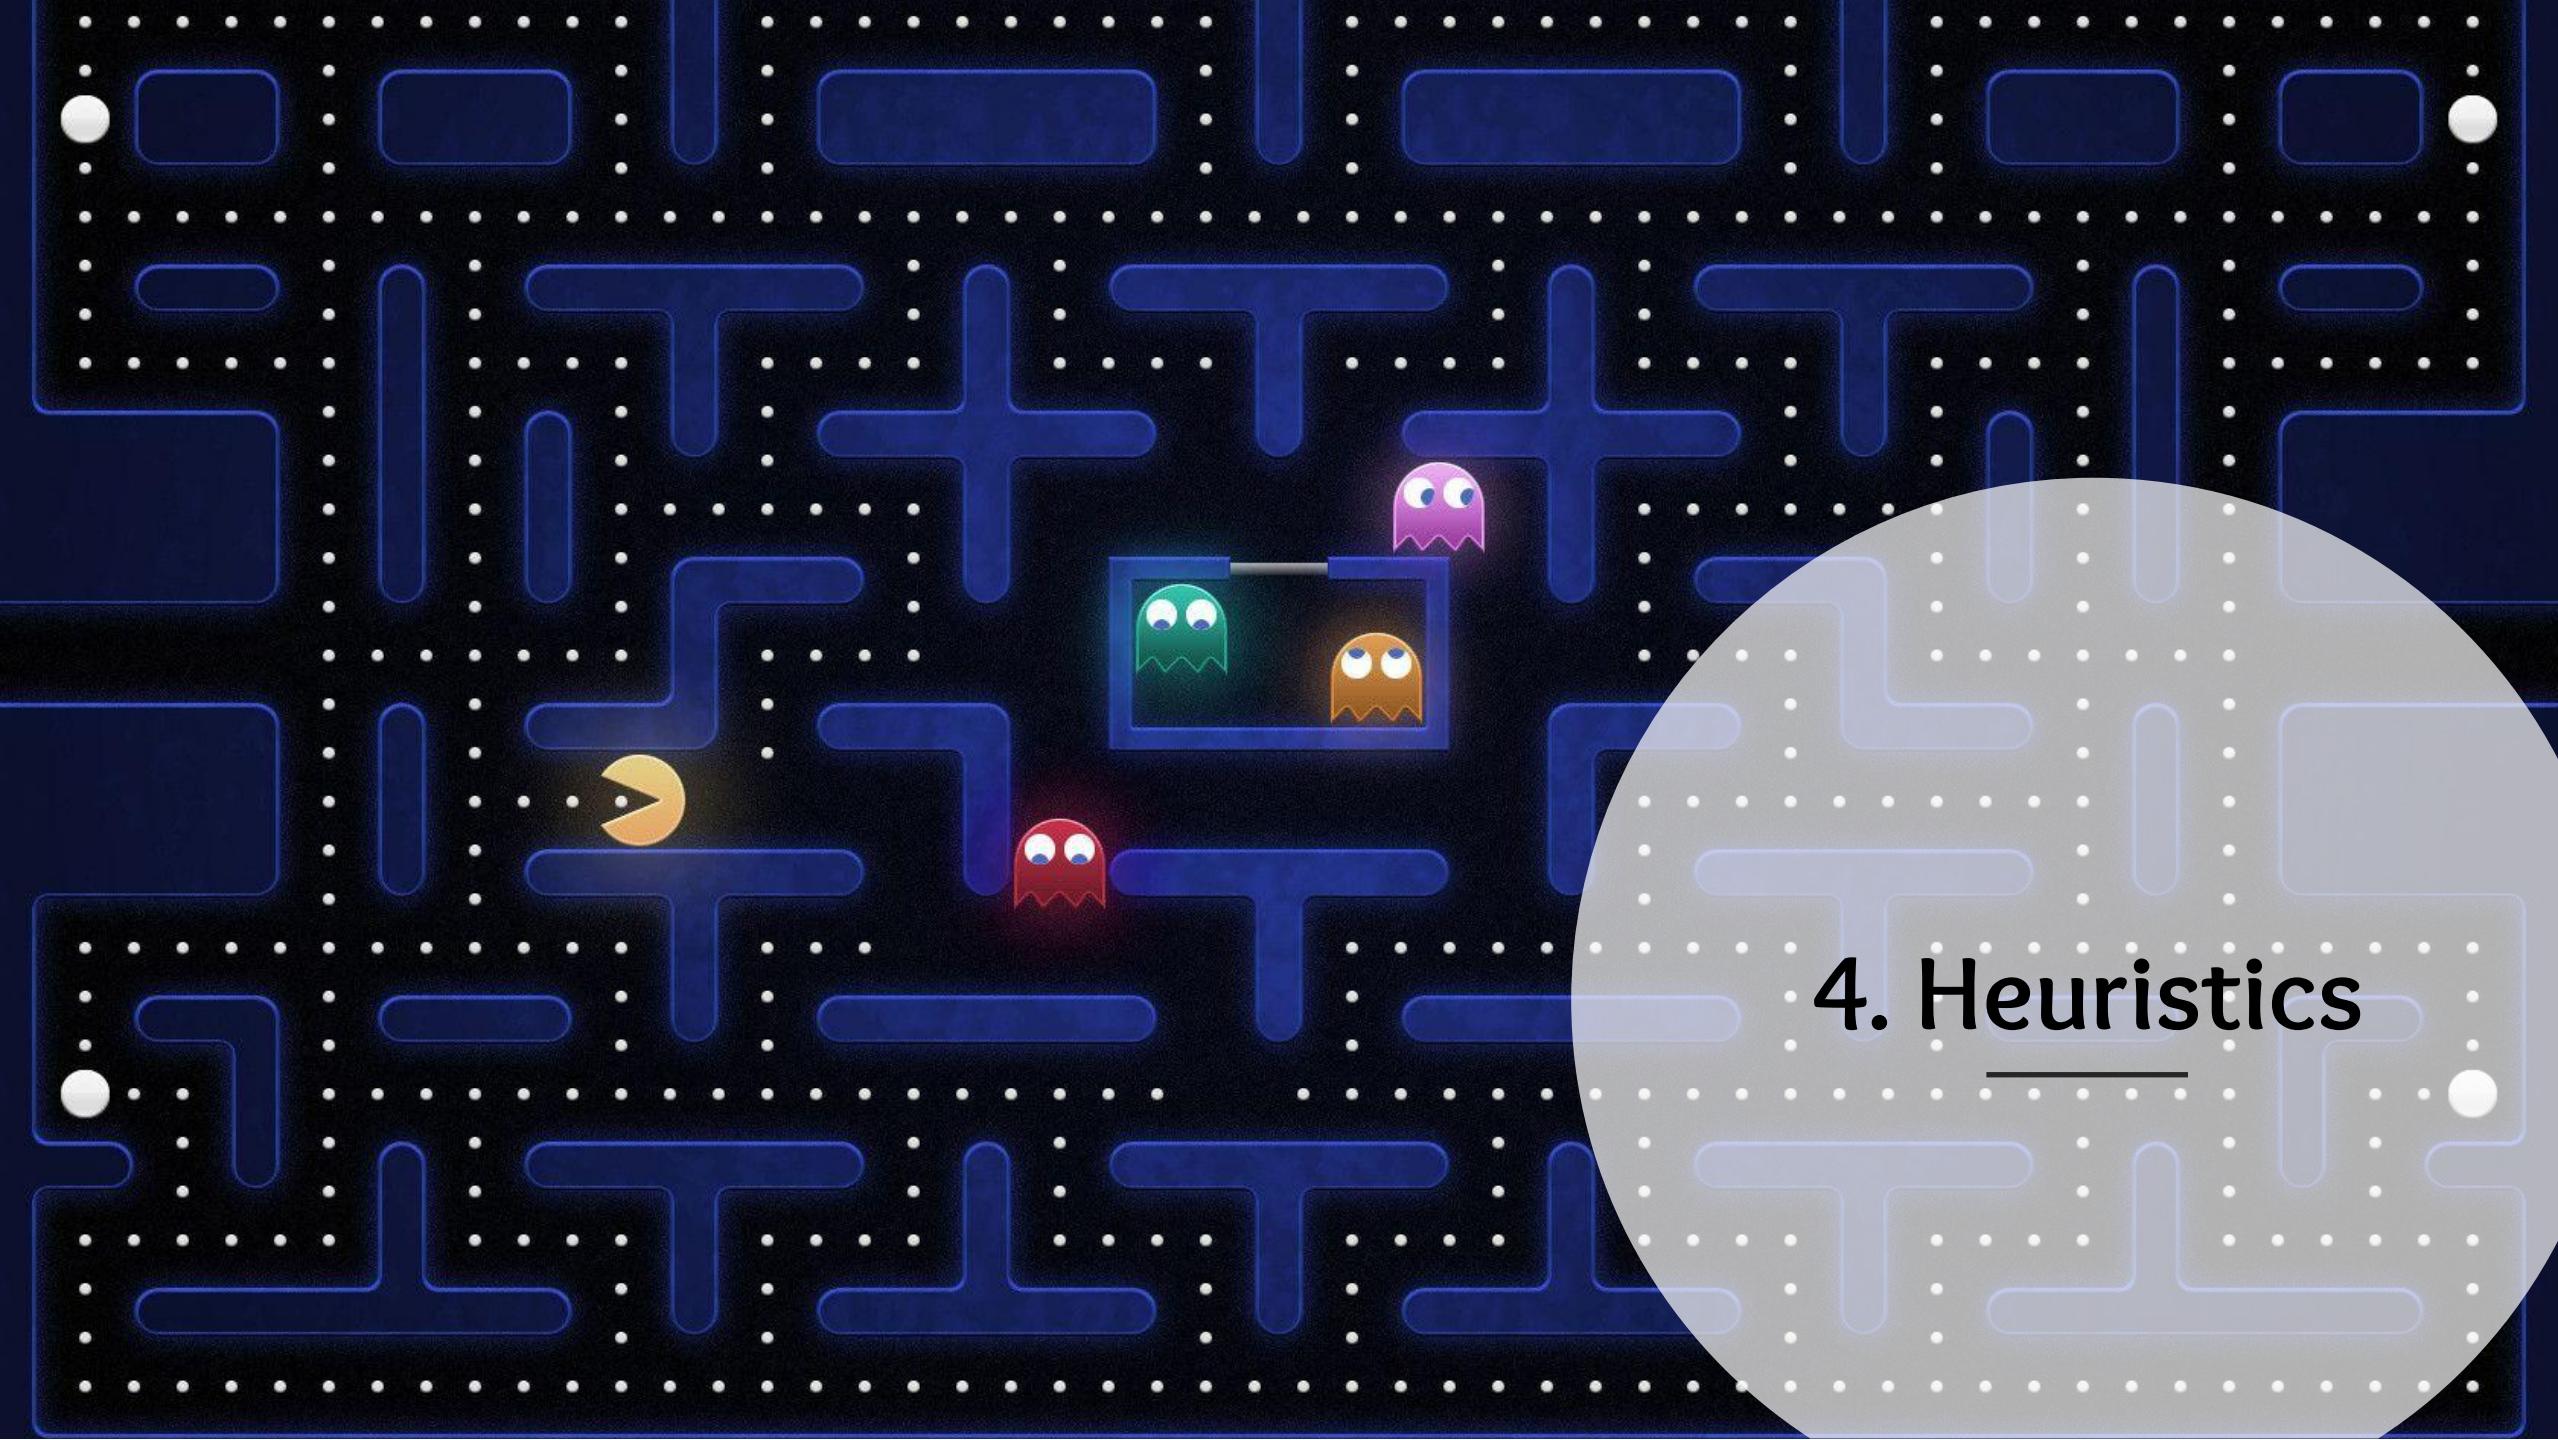
Goal Function: Efficiency Matters

- How often is our goal function called?
 - Often, every time we remove from our open list
- Don't make this function slow to compute

Modeling Corners Problem



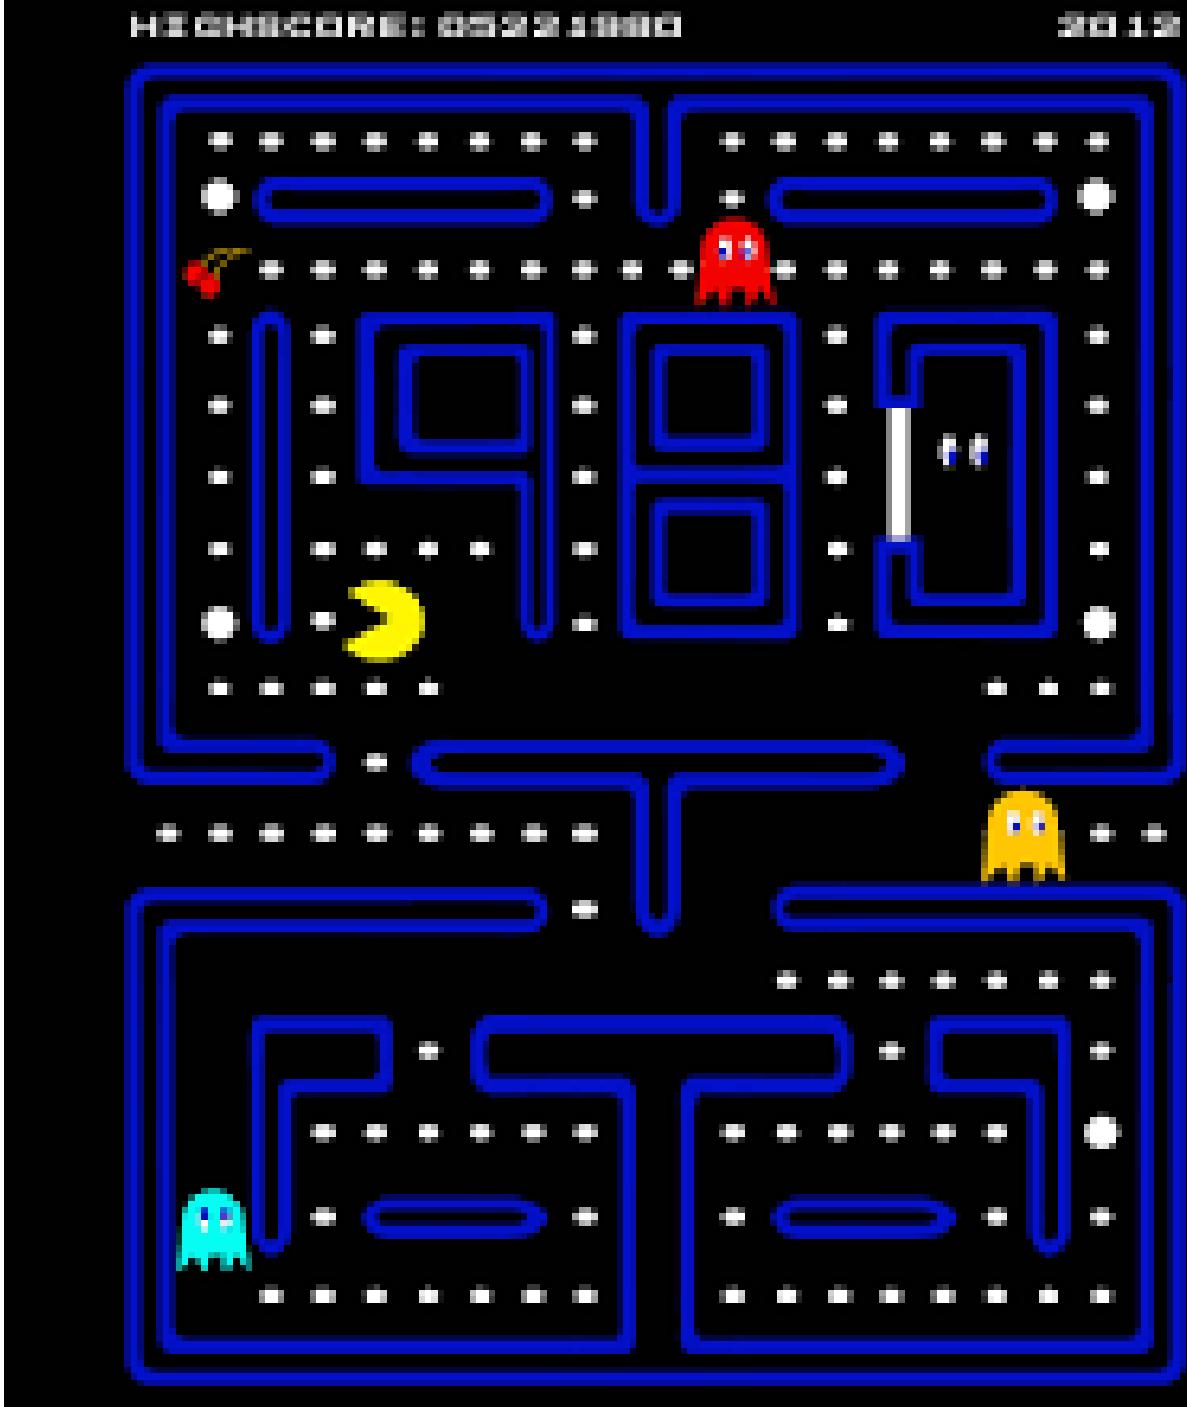
- States
 - Do we need to store the whole board?
 - Edit `getStartState()`
- Successor Function
 - Edit `getSuccessors(state)`
- Goal Function
 - Edit `isGoalState(state)`



4. Heuristics

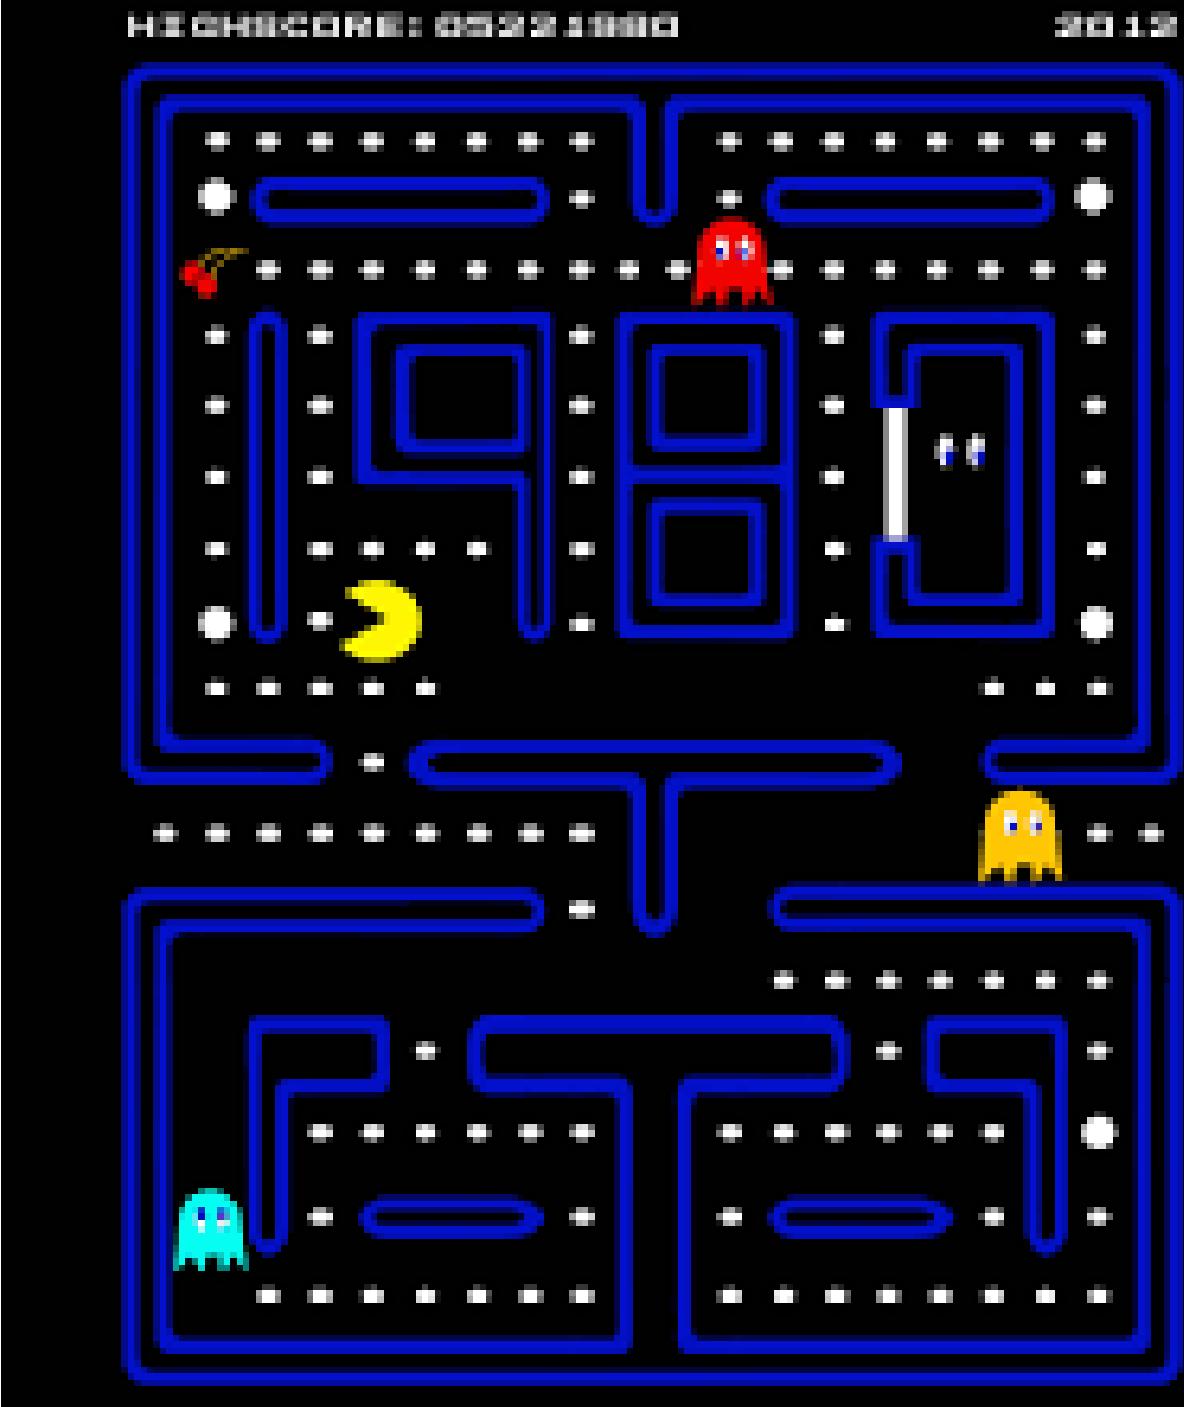
Heuristics: Definition!

- Quick Estimation / Rule of Thumb for how close you are to the goal.
- When you are AT the goal, the value of your heuristic should be **0**.
- Heuristics are **state specific!!**
- Give algorithms intuitions for focuses of optimization.



Heuristics: A* and UCS

- A* is designed to beat UCS because it is essentially UCS + heuristics.
- But only under certain conditions can A* beat UCS.



Heuristics: Good Heuristics

- Admissibility
- Consistency
- Also, a cute example to help y'all understand how heuristics are used in A*!



Admissible Heuristics

- **Never Overestimates!**
- Formal Definition:
 - Given a true heuristic represented as $h^*(s)$, then an admissible heuristic $h(s)$ satisfies:
$$h(s) \leq h^*(s) \text{ for all } s.$$
 - Admissibility ensures the optimality of A*.

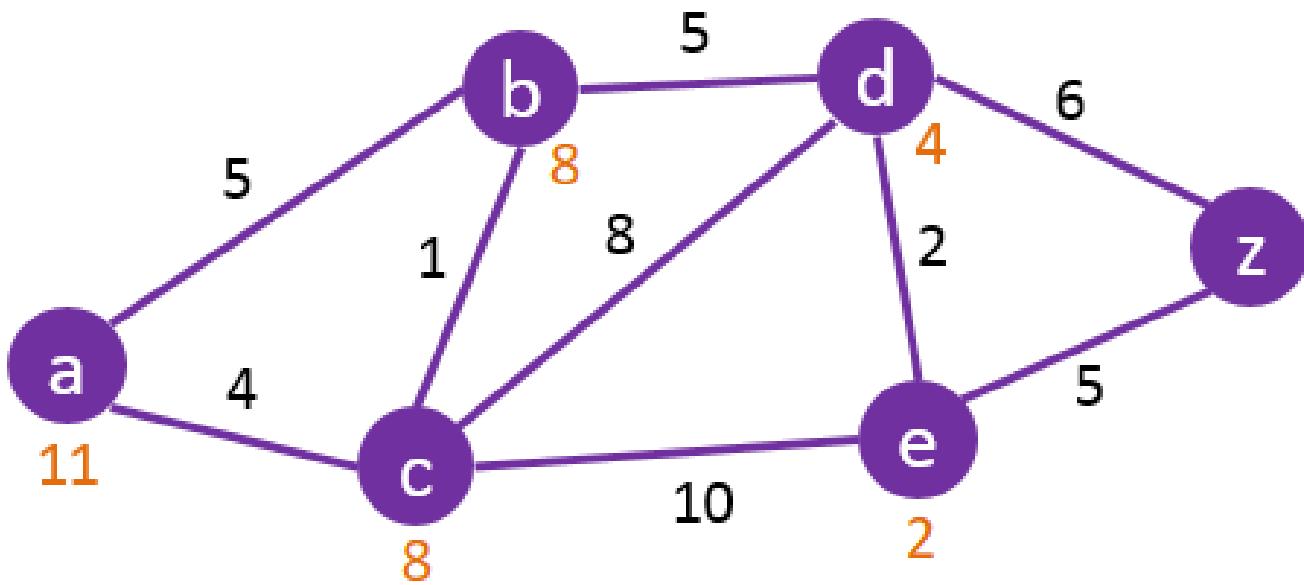


Consistent Heuristics

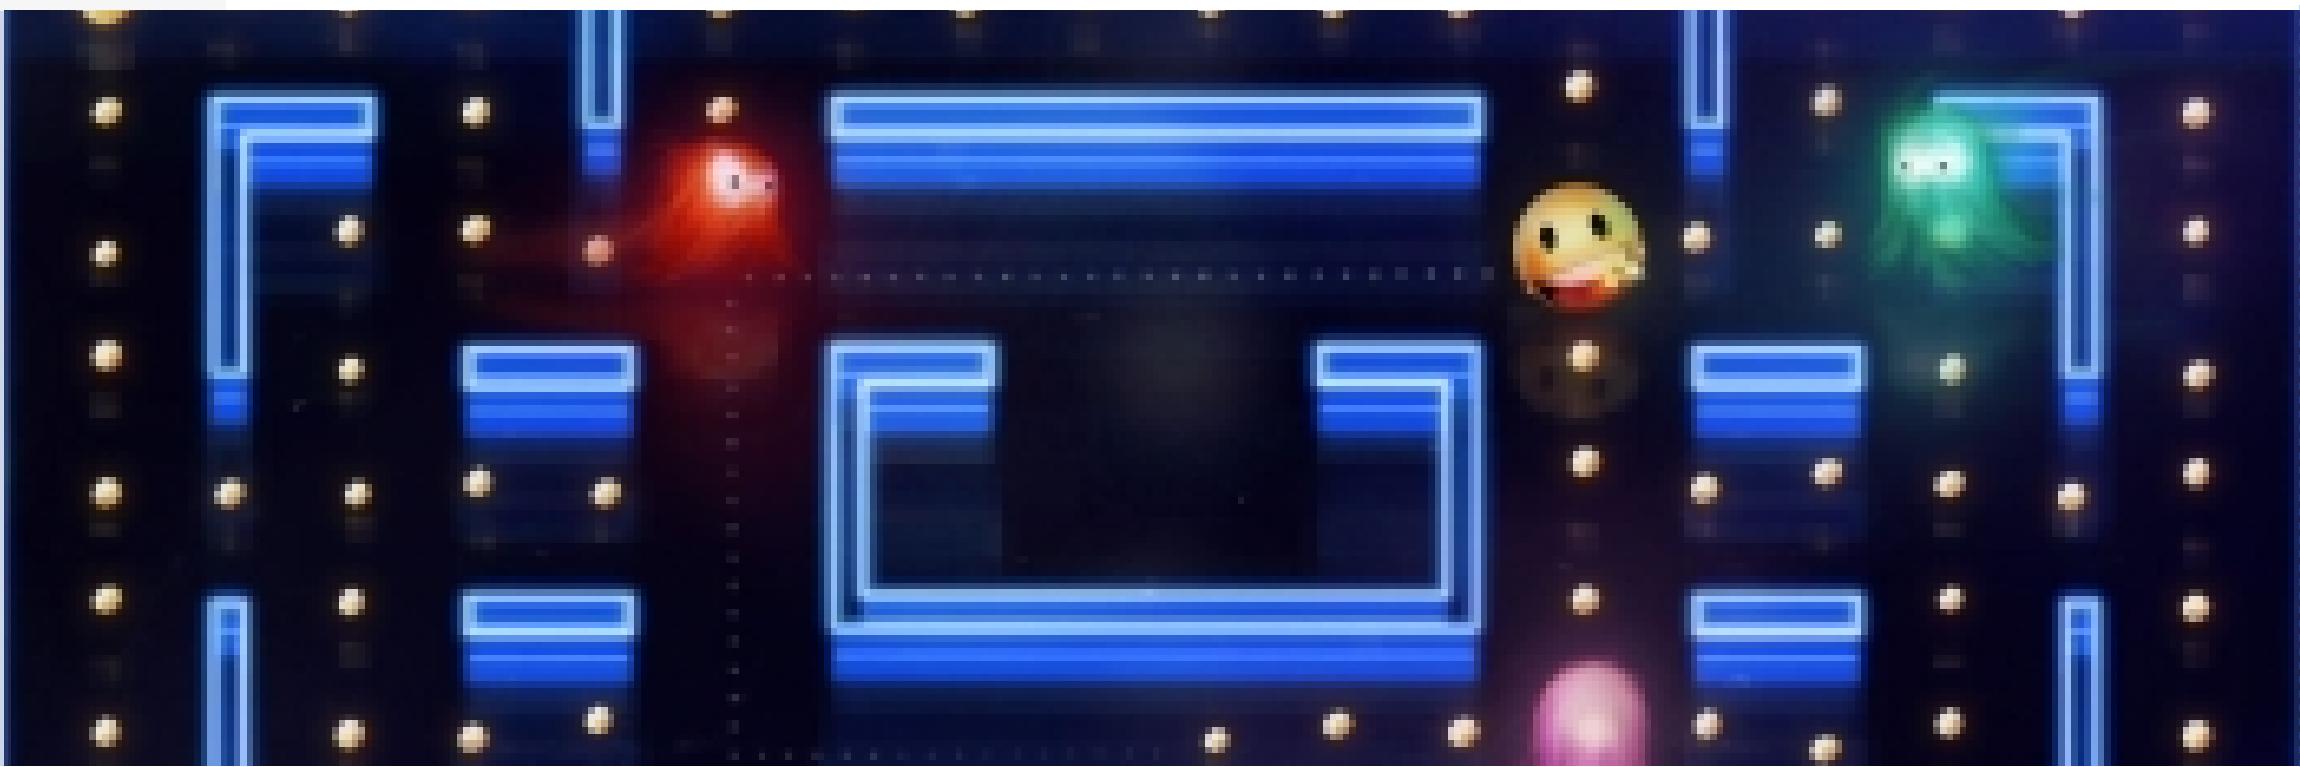
- The heuristic value should **decrease** when travelling from the parent to the child.
- $h(goal) = 0$
- $h(parent) - h(child) \leq \text{cost from parent to child}$



A* Example

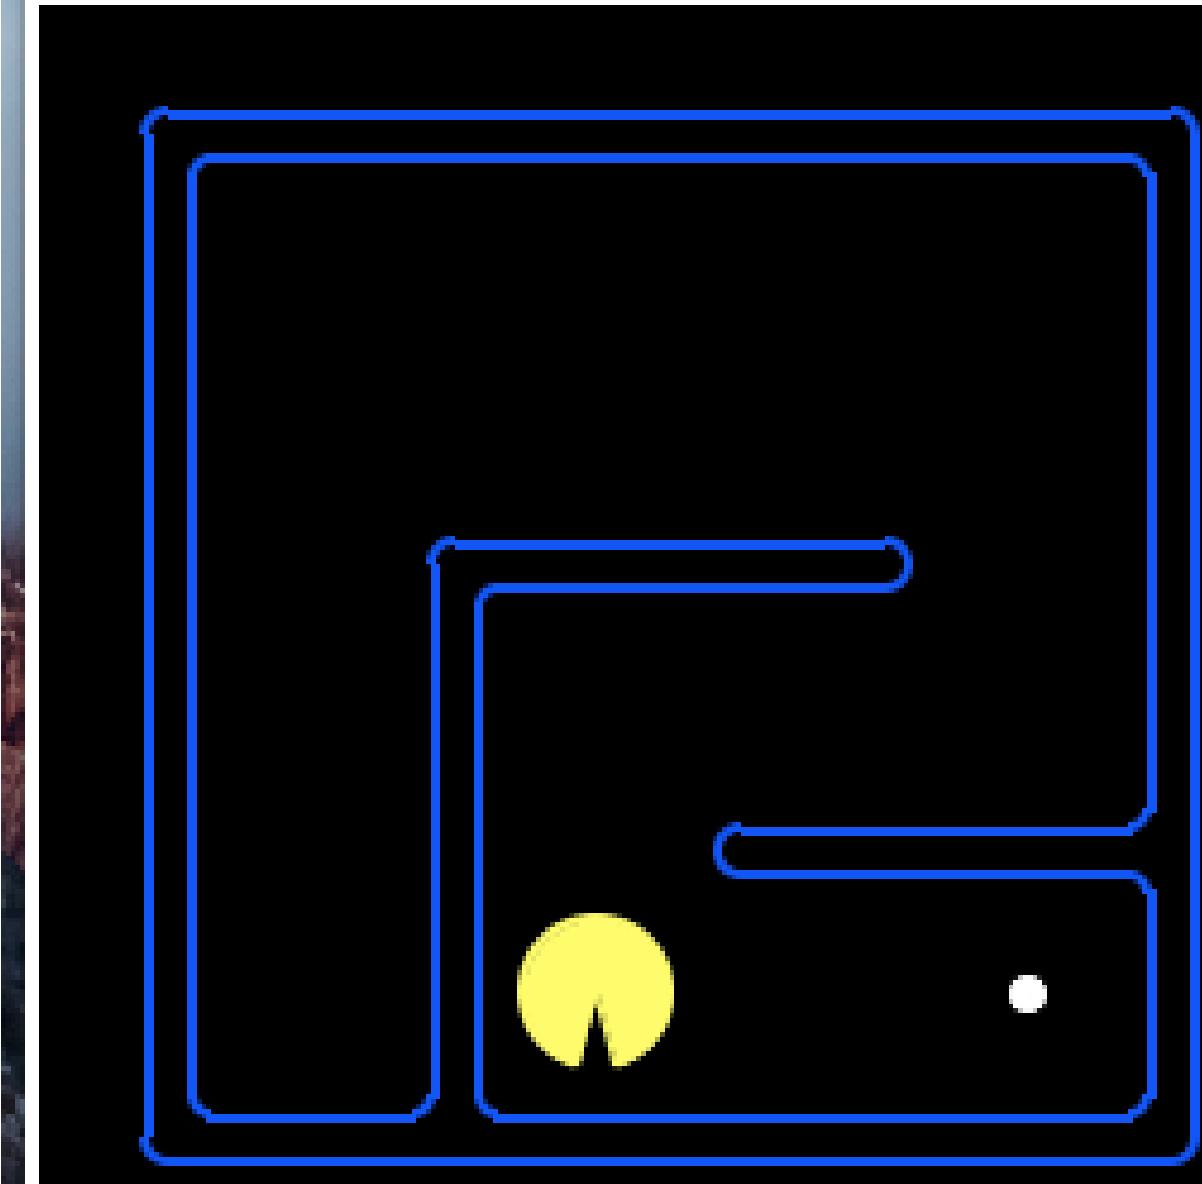


Corners Heuristics



Goal

- Traverse all four corners as fast as possible
- TO EAT FOOD PELLETS BECAUSE SUSTENANCE IS IMPORTANT
- You can define the states to allow you to keep track of corners you have visited to help you calculate the heuristics values

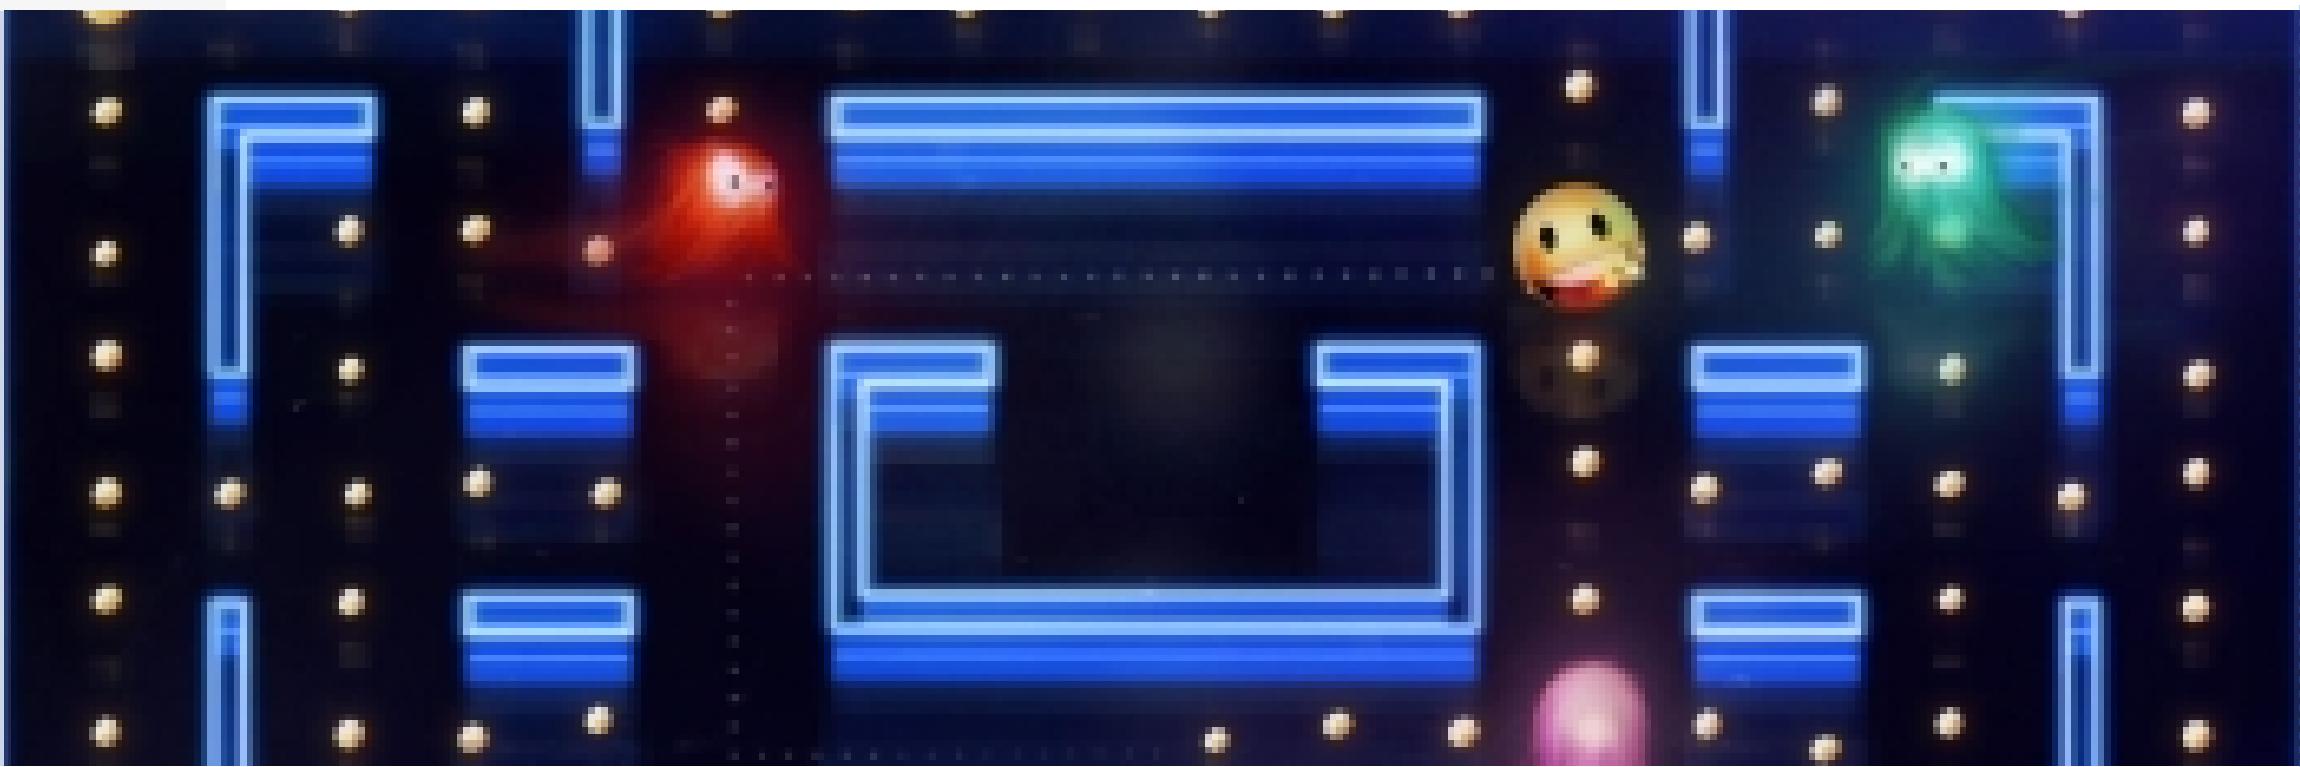


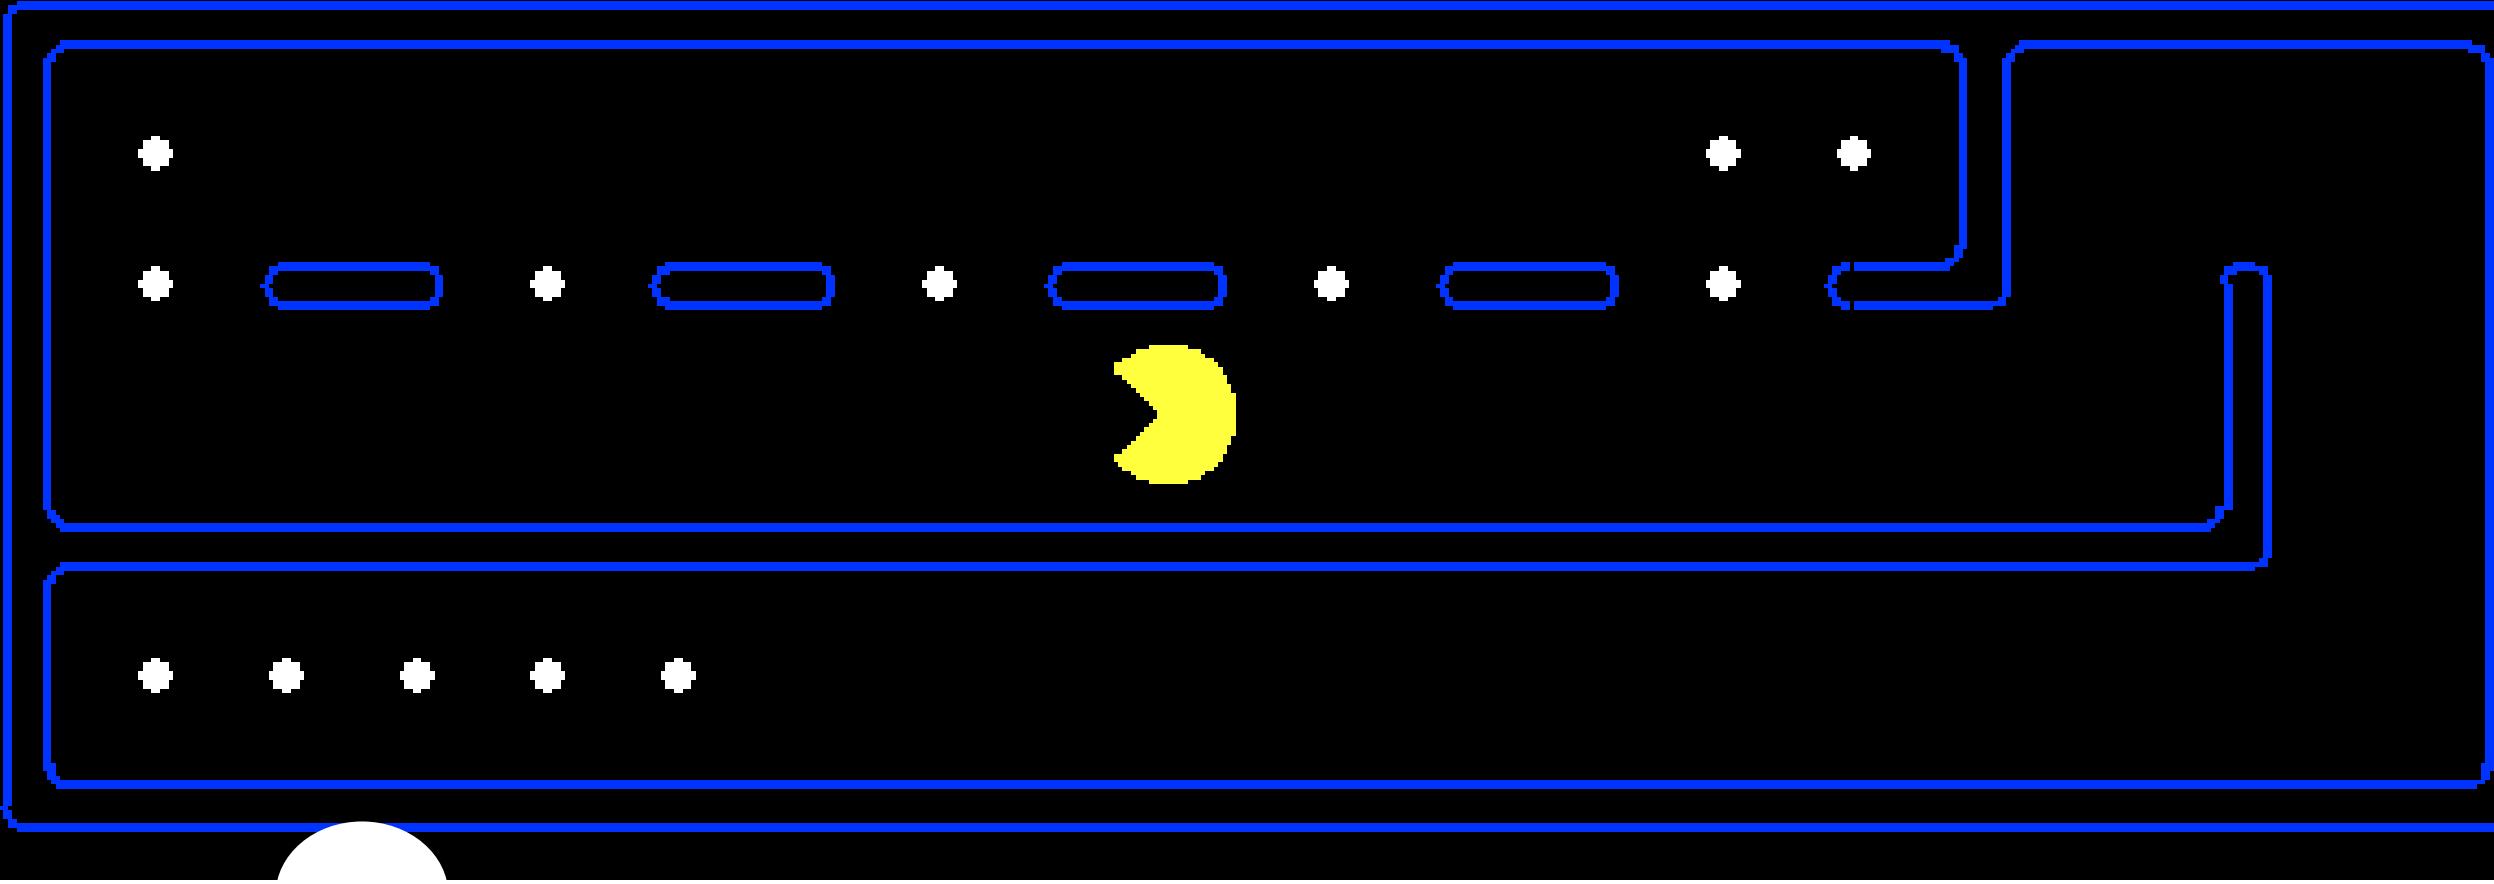
SCORE: 6

Some Common Approaches (and why they are wrong)

- Distance to the furthest pellet
- Distance to the closest pellet
 - Both approaches do not account for changing in the number of remaining pellets
- Sum of distances to all pellets
- Therefore, need to incorporate all possible information about all food pellets.

Food Heuristics





Hints

- Q6 is more statically defined (all pellets are placed at the corners)
- For Q7, the pellets are distributed across the whole grid
- Must come up with creative ways to incorporate all information



Questions?